

Chapter 10: Introduction to Artificial Neural Networks with Keras

- From Biological to Artificial Neurons
 - Biological Neurons
 - Logical Computations with Neurons
 - McCulloch and Pitts: network of neurons can compute any logical proposition
 - The Perceptron
 - Threshold logic unit, linear threshold unit
 - Weighted sum of inputs then step function
 - Step functions
 - Heaviside: 0/1
 - Sign: [-1,0,1]
 - Fully connected layer
 - Hebb's rule: cells that fire together, wire together
 - Perceptron learning rule:
 - $w_{i,j}^{(next\ step)} = w_{i,j} + \eta (y_j - \hat{y}_j) x_i$
 - η is the learning rate
 - Decision boundary of each output neuron is linear
 - Perceptrons incapable of learning complex patterns
 - Similar to Stochastic Gradient Descent
 - Limitations: unable to solve Exclusive OR (XOR) classification problem → MLPs can solve
 - The Multilayer Perceptron and Backpropagation
 - Regression MLPs
 - Classification MLPs
- Implementing MLPs with Keras
 - Installing TF2
 - Building an image classifier using the Sequential API
 - Sequential model: simplest model, single stack of layers connected sequentially
 - Layers:
 - Flatten: `X.reshape(-1,1)`
 - Dense: hidden layers, or output, initializes weights randomly
 - Can set `kernel_initializer` or `bias_initializer`
 - Compiling the model
 - `Sparse_categorical_crossentropy` for target class index
 - `Categorical_crossentropy` for one-hot encoded
 - Convert sparse labels (e.g., class indices) to one-hot vector labels use `keras.utils.to_categorical()`, to get back `np.argmax()`
 - `Validation_split` set instead of passing `validation_data`
 - `Class_weight` in `fit()` method to give larger weights to underrepresented classes
 - `Sample_weight` is per instance weights

- History.history contains loss and other metrics
 - Training error is running mean during epoch, validation error @ end of epoch
 - Not happy with model results:
 - Change learning rate then optimizer
 - Then # of layers, # of neurons, activation function @ each layer
 - Then batch size (default batch_size=32)
- Building a regression MLP using the Sequential API
 - keras.layers.Dense(1)
- Building complex models using the Functional API
 - Wide & Deep NN: connects all or part of inputs to output layer
 - Learns deep patterns and simple rules
 - Sequential forces all data to flow through full stack → simple patterns could get distorted
 - Called functional b/c calling layers like functions
 - Use case for multiple outputs:
 - Task driven: locate and classify main object in a picture
 - Multiple independent tasks based on same data
 - Regularization
- Using the Subclassing API to build dynamic models
 - Sequential and Functional are declarative → static
- Saving and restoring a model
 - Code: model.save('my_model.h5'),
keras.models.load_model('my_model.h5')
- Using callbacks
 - check_pt = keras.callbacks.ModelCheckpoint('my_model.h5')
 - model.fit(X_train, y_train, epochs = #, call_backs = [check_pt])
 - Early stopping:
 - early_stopping = keras.callbacks.EarlyStopping(patience=10, restore_best_weights=True)
 - model.fit(..., call_backs=[check_pt, early_stopping])
- Using TensorBoard for Visualization
 - Visualize learning curves
- Fine-tuning Neural Network Hyperparameters
 - Number of Hidden Layers
 - Parameter efficiency: deep networks can model more complex functions with fewer neurons than shallow ones
 - Learn hierarchical structure
 - Number of Neurons per Hidden Layer
 - Common to size to form a pyramid
 - Simpler to pick model with more layers and neurons than need and use early stopping
 - More bang from increasing layers than neurons

- Learning rate, batch size, and other hyperparameters
 - Learning rate: most important
 - If start small and increase, optimal rate just b4 loss function \uparrow
 - Optimizer: chap 11 discusses in detail
 - Batch size: affects performance and training time
 - GPU's can process efficiently \rightarrow training algo sees more instances per second, but may lead to instability, poor generalization
 - Activation function: ReLU usually good default
 - Number of iterations: don't need to tweak, stop early!