

Chapter 15: Processing Sequences Using RNNs and CNNs

- Introduction
 - Main difficulties
 - Unstable gradients: alleviated through dropout or layer normalization
 - Limited short-term memory: extended using LSTM or GRU
 - CNNs and regular NNs can also handle short sequences
- Recurrent Neurons and Layers
 - Intro
 - RNN connections forward and backward
 - Each time step each neuron receives both input vector x_t and output vector from previous time step, y_{t-1}
 - Each recurrent neuron has two sets of weights: inputs & t-1 output
 - RNN output: $y_{(t)} = \phi(w_x^T x_{(t)} + w_y^T y_{(t-1)} + b)$
 - Weights often concatenated vertically, x and y horizontally
 - Memory Cells
 - Since output of recurrent neuron (rn) is function of all inputs from previous steps → forms a kind of memory
 - Part of NN that preserves state across time steps is memory cell
 - Cell's state at time step t → h_t , hidden, is $f(h_{(t-1)}, x_t)$ or function of inputs and state at time step t-1
 - Basic cells: output = state, more complex not the same
 - Input and Output Sequences
 - Sequence-to-sequence network: feed previous stock prices, predict next price
 - Sequence-to-vector network: feed word sequence output sentiment score
 - Encoder-Decoder: S-2-v (encoder) → v-2-S (decoder), good for translation, better than word-by-word
- Training RNNs
 - Backpropagation through time
 - Forward pass through unrolled network
 - Output evaluated using cost function
 - Gradients flow backward through all outputs used by cost function
- Forecasting a Time series
 - Intro
 - Terms: univariate (one value per time step), multivariate (multi-values), imputation (postdicting, filling in missing values)
 - Note: input features represented as 3D array [batch_size, time_steps, dimensionality]
 - Baseline Metrics
 - Ways to compare model forecast
 - Naïve: predicts the last value
 - Fully connected network (FCN): add a flatten layer,
 - Implement a Simple RNN

- Single layer, single neuron, no need to specify length of input sequences, first input dimension = None
- SimpleRNN uses hyperbolic tangent, initial state (h_{init}) = 0,
 - Computes wt'd sum of values x_0 , applies $\tanh \rightarrow y_0$, which is new state h_0
 - New state, h_0 , passed to same rn with next input value x_1
 - Repeat until last time step
- Note: recurrent layers only return final output, to get them to return one output per time step \rightarrow return_sequences=True
- Parameters:
 - Linear: one per input and per time step + bias
 - Simple RNN: one per input and per hidden state dimension + bias
- Trend & seasonality
 - Time series models usually need to remove trend and seasonality through differencing, then train, then add back
 - RNNs generally don't have to remove
- Deep RNNS
 - Stack multiple layers of cells on one another
 - Note: set return_sequences=True for all recurrent layers but last
 - Why you don't want last layer to be SimpleRNN
 - Want prior hidden states to carry over into final output
 - SimpleRNN uses \tanh so predicted values need to be -1 to 1
 - Solution: replace with Dense layer
- Forecasting Several Time Steps Ahead
 - Predict values several steps ahead \rightarrow change target to desired time step
 - Predict sequences of steps ahead use first prediction as input for next
 - Errors \uparrow as number of time steps in sequence \uparrow
 - Predict next sequence of steps at every time step \rightarrow sequence-to-vector RNN becomes sequence-to-sequence RNN
 - More error gradients flowing through model
 - Flow through time, from output of each time step
 - Stabilize and speed up training [Why speed up?]
 - Causal model: only sees last time steps, but targets and values will overlap
 - To create sequence-to-sequence model \rightarrow return_sequences=True for all layers
 - Must apply Dense layer at every time step \rightarrow use TimeDistributed layer
 - Turns [batch_size, time_step, dim] into [batch_size * time_step, dim] and back
 - All outputs needed during training, only last output at last time step useful for predictions and evaluation
 - Note: error bars useful to include in predictions. Use MCDropout within each cell. After training forecast many times and compute mu and sigma of predictions

- Handling Long Sequences
 - Intro
 - Training RNN on long sequences means many time steps, means deep NN
→ unstable gradients problem
 - RNN will forget earlier inputs in long sequence
 - Fighting the Unstable Gradients Problem
 - Common solutions: good parameter initialization, faster optimizers, dropout
 - Non-saturating activation not as successful
 - Early wts used later → may increase or explode
 - Other methods: reduce learning rate, using saturating activation function, gradient clipping
 - Batch normalization not as efficient: can't use between time steps, only between layers, improves only slightly if applied initially
 - Reason: BN layer used at each time step with same parameters, regardless of actual scale and offset of inputs and hidden state
 - Layer Normalization works better: normalizes across feature dimensions
 - Can compute stats on fly at each time step independently of each instance
 - Behaves same during training and testing, does not need to use EMA to estimate feature stats across instances
 - LN learns a scale and offset parameter for each input, used after linear combination of inputs and hidden states
 - All Keras layers (except .RNN) have dropout and recurrent_dropout hyperparameters, apply to inputs or to hidden state
 - Tackling the Short-Term Memory Problem
 - Intro
 - Information lost at each step when traversing an RNN
 - Solution: cells with long-term memory
 - LSTM cells
 - Apply with `keras.layers.LSTM` or with `layers.RNN(keras.layers.LSTMCell())`
 - Main idea: learn what to store in long term state, what to throw away, and what to read from it
 - Each time step: some memories dropped, some added, long-term state copied passed, activated, and filtered to output → short term state (h_t)
 - LSTM learns to recognize important input, store it in long-term state, preserve for as long as needed, and then extract when needed
 - Peephole connections: look at long-term state, only in experimental version
 - GRU cells
 - Gated Recurrent Unit: simplified version of LSTM, performs just as well
 - State vectors merged into h , single gate controller, no output gate

- New gate controller controls which part of previous state shown to main layer
- Issues
 - LSTM and GRUs still have problem handling sequences longer than 100 steps
- Other structures
 - Using 1D convolutional layers to process sequences
 - 1D convolutional layer slides several kernels across a sequence producing 1D feature map per kernel
 - Wavenet
 - Stacks 1D convolutional layers, doubling dilation rate at every layer
 - Dilation rate: how spread apart neuron's inputs are
 - First conv layer sees 2 time steps, next 4, next 8, up to 512